

Hidden Markov Models for Protein Dynamics

September 24, 2013

1 Motivation

Traditionally, we have clustered conformation-space into many discrete voronoi cells and used this clustering to build a large transition matrix. This results in implied timescales that are consistently faster than they should be.

The motivation of this work can be found in considering a simple, two-well, 1D potential. What is the optimal clustering? There are only two physically stable states, but if we used two *hard* states, conformations on the cusp of the barrier are considered to be the same as conformations at the bottom of the well. If the particle moving under this potential diffuses for a bit at the top of the potential, crossing the line that distinguishes our two hard states, it generates many counts, distorting the dynamics.

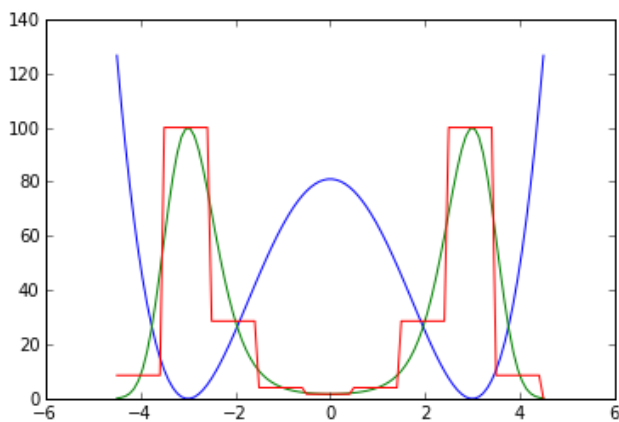


Figure 1: A two-well potential (blue) and its stationary distribution (green). A traditional MSM will produce a discretized approximation (red) to the true stationary distribution.

The theoretical underpinnings of Markov State Models is that they are used to estimate the propagator. The true propagator is continuous in \mathbb{R}^{6N} , and the typical approach is to use indicator basis functions to approximate it. This gives step-like eigenfunctions whose resolution increases with increasing number of states—see figure 1, red line. One can imagine using a different basis set instead. Two Gaussian functions, for example, could describe dynamics on this potential much better than even a high number of step functions with many fewer parameters.

2 HMMs and Baum-Welch

Deriving a transition matrix from non-overlapping hard states is trivial: for lag time τ , if the conformation moves from state i to state j , we increment the entry in the count matrix, C_{ij} . The transition matrix is a row-stochastic normalization of this count matrix.

It is impossible to define counts in this way for *fuzzy* states. One might imagine describing a point at time t not simply as belonging to state i , but rather as having a vector of memberships \mathbf{m} in each fuzzy state. One might think that you can get the transition matrix by summing the outer product between each time-pair of membership vectors as in REF

$$\sum |\mathbf{m}_t\rangle \langle \mathbf{m}_{t+1}| \quad (1)$$

This works for the case of quantized memberships, but creates arbitrarily small transition probabilities due to the multiplication of many numbers that are less than one. In fact, the problem of coming up with a transition matrix from fuzzy states is non-trivial. Going back to our two well potential, consider a conformation sitting at the top of the barrier. It has membership vector $\mathbf{m}_0 = (0.5, 0.5)$. After one time step, it is in the same place, $\mathbf{m}_1 = (0.5, 0.5)$. What is the proper transition matrix? Is it the identity matrix? Is it $\frac{1}{2}$ in each position? It's difficult to say.

We require a learning algorithm that can infer from entire trajectories, probable state definitions and the paths through those states.

Instead of treating the system as a Markov model, we can consider it a Hidden Markov Model (HMM). The conformation at a particular time is the observed quantity, or *emission*, and the small number of (potentially metastable) states are the hidden states of the model.

The Baum-Welch algorithm is a particular expectation-maximization (EM) algorithm to simultaneously learn the parameters of the hidden states ('emission distributions') as well as the rates of transition through those states (the hidden transition matrix)

3 Methods

The Baum-Welch algorithm can only find local maxima in likelihood. It is important to seed the algorithm with initial transition matrices and state definitions that are close to the actual ones. In our approach, we use standard gaussian mixture modelling from the scikit-learn package to define initial gaussian emission distributions. From this density based approach, we can calculate the BIC of the mixture model to determine the best number of states. Each conformation is assigned membership values (zero to one) in each gaussian state. As a first approximation for calculating the transition matrix, conformations are assigned to the state in which their membership is highest. A normal (relatively few-state) MSM is built from these assignments, and this transition matrix is used to seed the Baum-Welch algorithm.

After the initial guesses for state definition and transition matrix have been made, the algorithm, as implemented in the General HMM (ghmm) package for C/python is run on a set of trajectories. Optimized parameters for the emission distributions and hidden transition matrix are returned.

There is probably some math to extract the eigenvector/values from this hidden transition matrix.

$$T(\tau) = Q^T \tilde{T} Q \quad (2a)$$

$$T(2\tau) = Q^T \tilde{T}^2 Q \quad (2b)$$

$$\neq Q^T \tilde{T} Q Q^T \tilde{T} Q = T(\tau)^2 \quad (2c)$$

4 The Muller Potential

4.1 Qualitative

As validation, the HMM procedure was run on data for the Muller Potential. A high-resolution (32,600 state) transition matrix was built based on the potential and an arbitrary ‘temperature’. Truly markovian trajectories can be generated from this large transition matrix.

In this example, we allow unrestricted covariance matrices for each of the three states.

2 shows the sample trajectories generated from the analytical transition matrix. First, a mixture model was generated (figure 3). Using the Baum-Welch algorithm, the emission distributions (figure 4) and hidden transition matrix were optimized. We visualize the first three eigenfunctions (figure 5) for the generating transition matrix (first column) vs. those derived from the HMM (second column).

A given The HMM eigenfunction i was plotted as $\mathbf{O}^T \cdot \pi_i$, where \mathbf{O} is a vector of the emission probabilities and π_i is the i -th eigenvector of the hidden transition matrix. This may or may not be valid.

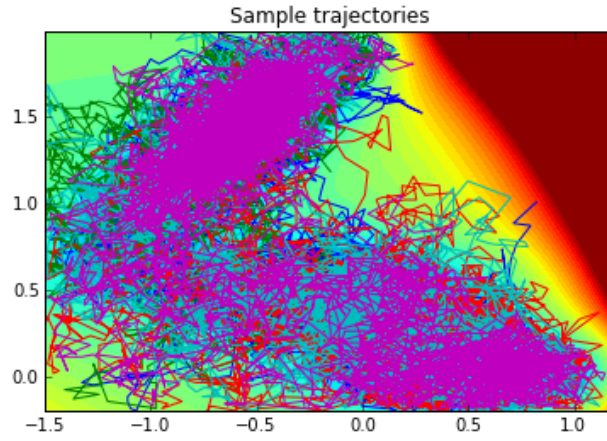


Figure 2: 5 trajectories of length 250,000 steps

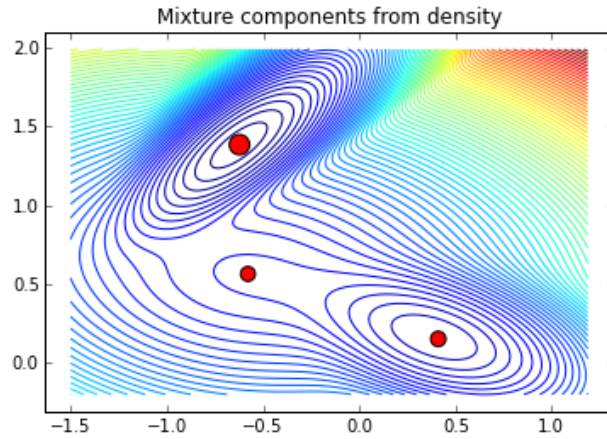


Figure 3: The initial mixture model without regard for the time-series nature of the data. The model reflects the lack of density in the transition region.

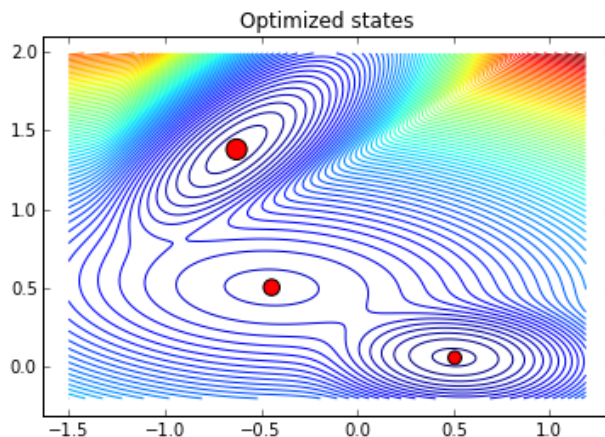


Figure 4: The optimized hidden states after running the Baum-Welch algorithm. Note that this is a sum of three distinct probability distributions, each normalized to one.

4.2 Quantitative Validation

We hypothesized that HMMs could more accurately predict the implied timescales (proportional to the eigenvalues of the transition matrix) than a normal MSM all else being equal. A validation script was run (run time 2 days) that performed HMM and MSM analysis doing all possible permutations of the following parameters:

- number of clusters(k): 2 – 19
- lag time: 1 – 29, every two
- random seed (repeat each case): 10 of them

The models were trained on 5 trajectories of length 100,000 steps, strided by 10. For the future, it would be beneficial to re-do this analysis without striding the data.

4.3 Right now

Re-running with 10 trajectories of 500,000 steps with stride of 5.

Further thoughts: Comparing number of clusters to number of clusters is kind of unfair. We should use the number of parameters. In MSM, the number of parameters is just $dim * k$. In the muller potential it is $(dim^2 + dim) * k$. The

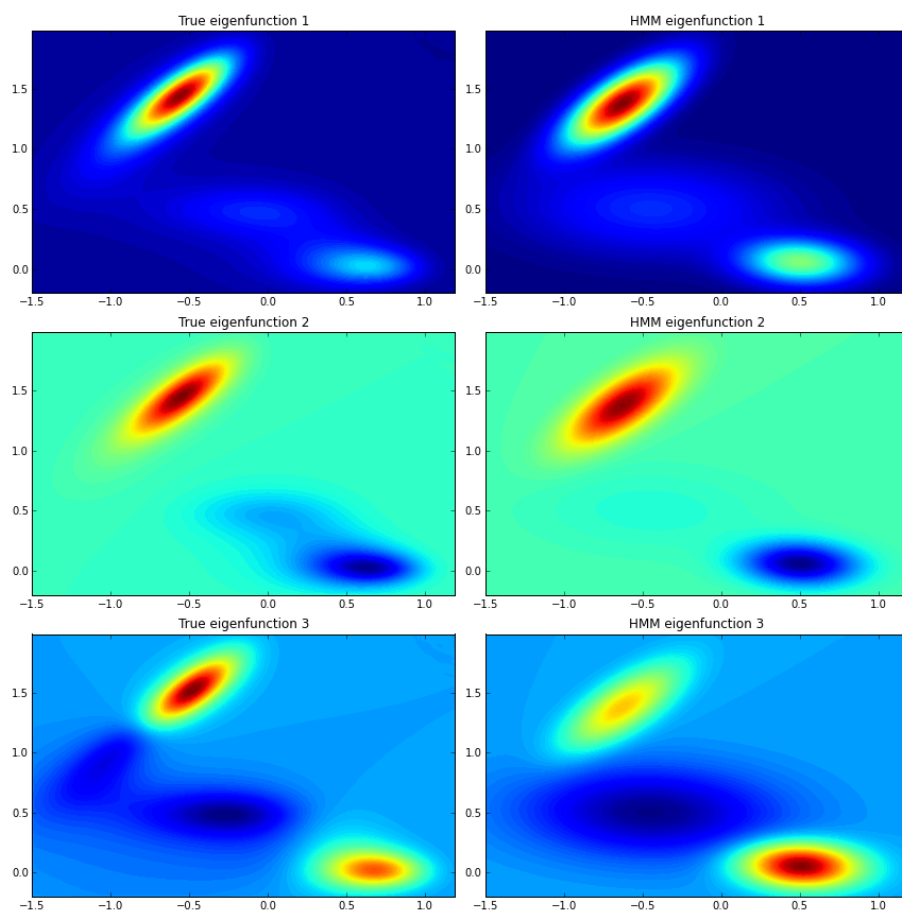


Figure 5: Visualization of eigenfunctions

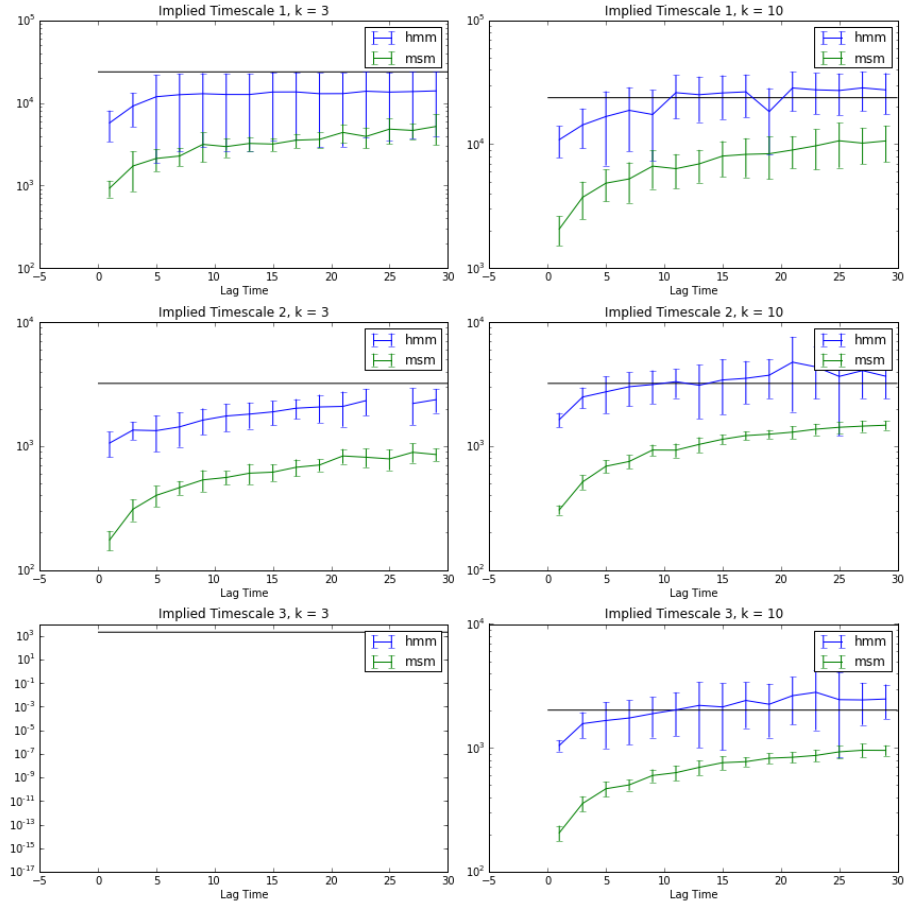


Figure 6: Implied timescales vs. lag time.

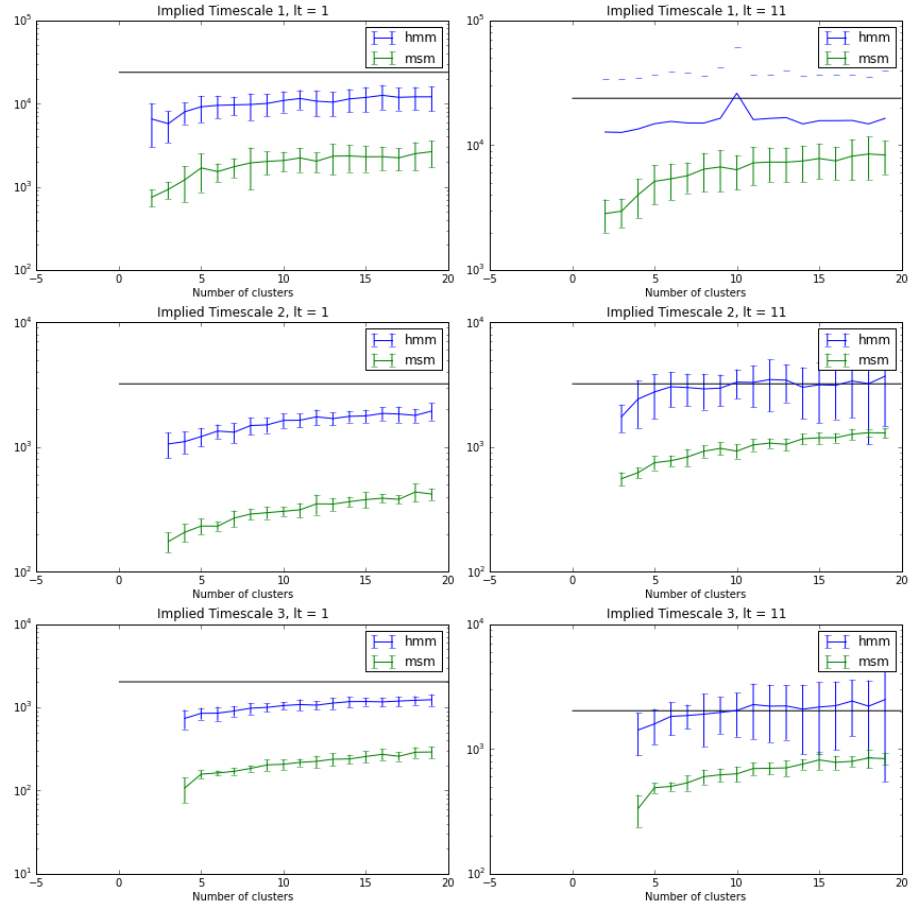


Figure 7: Implied timescales vs. number of clusters

HMM uses a factor of $dim + 1$ more parameters than msm for a given number of clusters Using $dim=2$, $k_{msm} = 3 * k_{hmm}$. I'm doing this in the new run.

5 Alanine Dipeptide

We have applied the HMM model to alanine dipeptide. It is acting fishy. I would like a better way to visualize the resulting emission distributions; maybe as ellipses. It is possible that there is one that is just too wide.

The HMM requires that conformations be vectorizable. This is hard to do. I'm suspecting that using sines and cosines is still messing things up. Those values are confined within -1 to 1, whereas the emissions use non-wrapping real numbers.

I'm going to try using atom-pair distances.

6 Other thoughts

Perhaps inject some noise into the baum welch to find a better maximum.

Using R is probably of little benefit.